



Prácticas de Administración de Bases de Datos

Grado en Ingeniería Informática

PRÁCTICA 5

Procesamiento y optimización de consultas

Ejercicios

1. Tras una actualización de la base de datos (desde la versión 10.2.0.4), se ha informado de que algunas consultas no parecen ser tan eficientes como anteriormente. Se le pide al administrador que mientras no se encuentren posibles optimizaciones, cambie algún parámetro del optimizador de consultas para que funcione igual que en la versión anterior. ¿Cómo podría realizarse?

```
ALTER SYSTEM SET optimizer_features_enable='10.2.0.4';
```

```
SHOW PARAMETER optimizer_features_enable
```

2. ¿Qué tendría que hacerse para que el optimizador no usara estadísticas dinámicas en la sesión actual?

```
ALTER SESSION SET optimizer_dynamic_sampling=0;
```

3. Un proceso realiza consultas a la base de datos y requiere bajos tiempos de respuesta, aunque no necesita todos los datos en un primer momento. ¿Cómo podría cambiarse el comportamiento del optimizador para conseguirlo?

```
ALTER SESSION SET OPTIMIZER_MODE=FIRST_ROWS_1;
```

4. Examina el contenido del script *catplan.sql*.

5. Lanza EXPLAIN PLAN (guardándolo con el `statement_id` *sentencia1*) para la siguiente consulta:

```
SELECT phone_number
  FROM hr.employees
 WHERE phone_number LIKE '%16%';
```

Muestra el plan mediante UTLXPLS.SQL y mediante DBMS_XPLAN.DISPLAY.
¿Qué diferencias hay entre hacerlo de una u otra manera?

Muestra el plan con el mayor nivel de detalle y también con el menor nivel de detalle.

```
EXPLAIN PLAN SET statement_id = 'sentencia1' FOR SELECT phone_number
FROM hr.employees
WHERE phone_number LIKE '%16%';
```

//La ruta puede cambiar dependiendo de nuestra instalación

```
@C:\app\usuario\product\11.2.0\dbhome_1\RDBMS\ADMIN\UTLXPLS.SQL
```

```
SELECT PLAN_TABLE_OUTPUT
FROM TABLE(DBMS_XPLAN.DISPLAY(null, 'sentencia1'));
```

```
SELECT PLAN_TABLE_OUTPUT
FROM TABLE(DBMS_XPLAN.DISPLAY(null, 'sentencia1', 'ALL'));
```

```
SELECT PLAN_TABLE_OUTPUT
FROM TABLE(DBMS_XPLAN.DISPLAY(null, 'sentencia1', 'BASIC'));
```

6. Lanza EXPLAIN PLAN para la siguiente consulta:

```
SELECT *
FROM hr.employees
ORDER BY last_name;
```

¿Cambia algo del plan si modificas para la actual sesión el valor del parámetro DB_FILE_MULTIBLOCK_READ_COUNT a 1? ¿Por qué?.

```
SHOW PARAMETER DB_FILE_MULTIBLOCK_READ_COUNT
```

```
EXPLAIN PLAN FOR SELECT *
FROM hr.employees
ORDER BY last_name;
```

```
SELECT PLAN_TABLE_OUTPUT
FROM TABLE(DBMS_XPLAN.DISPLAY());
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		107	7383	4 (25)	00:00:01
1	SORT ORDER BY		107	7383	4 (25)	00:00:01
2	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01

```
ALTER SESSION SET DB_FILE_MULTIBLOCK_READ_COUNT=1;
```

--Ahora el plan de ejecución queda--

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		107	7383	7 (15)	00:00:01
1	SORT ORDER BY		107	7383	7 (15)	00:00:01
2	TABLE ACCESS FULL	EMPLOYEES	107	7383	6 (0)	00:00:01

```
ALTER SESSION SET DB_FILE_MULTIBLOCK_READ_COUNT=128;
```

7. Lanza EXPLAIN PLAN para la siguiente consulta:

```
SELECT *
FROM sh.customers
WHERE cust_city='Los Angeles'
AND cust_state_province='CA';
```

Examina el plan resultante. ¿Cambia algo del plan si modificas para la actual sesión el valor del parámetro OPTIMIZER_DYNAMIC_SAMPLING a 10?.

```
EXPLAIN PLAN FOR
SELECT *
FROM sh.customers
WHERE cust_city='Los Angeles'
AND cust_state_province='CA';
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	181	261 (1)	00:00:04
* 1	TABLE ACCESS FULL	CUSTOMERS	1	181	261 (1)	00:00:04

Predicate Information (identified by operation id):

```
1 - filter("CUST_CITY"='Los Angeles' AND "CUST_STATE_PROVINCE"='CA')
```

ALTER SESSION SET OPTIMIZER_DYNAMIC_SAMPLING=10;

--ahora el plan queda...--

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		932	164K	261 (1)	00:00:04
* 1	TABLE ACCESS FULL	CUSTOMERS	932	164K	261 (1)	00:00:04

Predicate Information (identified by operation id):

```
1 - filter("CUST_CITY"='Los Angeles' AND "CUST_STATE_PROVINCE"='CA')
```

Note

```
- dynamic sampling used for this statement (level=10)
```

Ha actualizado estadísticas

8. En cierto código se ha encontrado la siguiente consulta:

```
SELECT /*+ FULL(e) */ employee_id, last_name
FROM hr.employees e
WHERE last_name LIKE '%Smith%';
```

Analiza su plan de ejecución actual, si el plan de ejecución hace uso de algún índice y cómo podría mejorarse.

EXPLAIN PLAN FOR
SELECT /*+ FULL(e) */ employee_id, last_name
FROM hr.employees e
WHERE last_name LIKE '%Smith%';

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		5	60	3 (0)	00:00:01
* 1	TABLE ACCESS FULL	EMPLOYEES	5	60	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
1 - filter("LAST_NAME" LIKE '%Smith%')
```

EXPLAIN PLAN FOR
SELECT employee_id, last_name
FROM hr.employees e
WHERE last_name LIKE '%Smith%';

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		5	60	3 (34)	00:00:01
1	VIEW JOIN	index\$_join\$_001	5	60	3 (34)	00:00:01

*	2		HASH											
	3		INDEX FAST FULL SCAN		EMP_EMP_ID_PK		5		60		1	(0)		00:00:01
*	4		INDEX FAST FULL SCAN		EMP_NAME_IX		5		60		1	(0)		00:00:01

Predicate Information (identified by operation id):

```

2 - access(ROWID=ROWID)
4 - filter("LAST_NAME" LIKE '%Smith%')

```

EXPLAIN PLAN FOR
SELECT /*+ INDEX(e EMP_NAME_IX) */ employee_id, last_name FROM hr.employees e
WHERE last_name LIKE '%Smith%';

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		5	60	2 (0)	00:00:01	
1	TABLE ACCESS BY INDEX ROWID	EMPLOYEES	5	60	2 (0)	00:00:01	
*	2	INDEX FULL SCAN	EMP_NAME_IX	5		1 (0)	00:00:01

9. En cierto código se ha encontrado la siguiente consulta:

```

SELECT /*+ NO_USE_HASH(o l) */
       o.customer_id, l.unit_price * l.quantity
FROM   oe.orders o, oe.order_items l
WHERE  l.order_id = o.order_id;

```

Analiza su plan de ejecución y estudia si podría mejorarse.

EXPLAIN PLAN FOR
SELECT /*+ NO_USE_HASH(o l) */
o.customer_id, l.unit_price * l.quantity FROM oe.orders o, oe.order_items l
WHERE l.order_id = o.order_id;

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		665	13300	8 (38)	00:00:01	
1	MERGE JOIN		665	13300	8 (38)	00:00:01	
2	SORT JOIN		105	840	4 (50)	00:00:01	
3	VIEW	index\$_join\$_001	105	840	3 (34)	00:00:01	
*	4	HASH JOIN					
5	INDEX FAST FULL SCAN	ORDER_PK	105	840	1 (0)	00:00:01	
6	INDEX FAST FULL SCAN	ORD_CUSTOMER_IX	105	840	1 (0)	00:00:01	
*	7	SORT JOIN		665	7980	4 (25)	00:00:01
8	TABLE ACCESS FULL	ORDER_ITEMS	665	7980	3 (0)	00:00:01	

Predicate Information (identified by operation id):

```

4 - access(ROWID=ROWID)
7 - access("L"."ORDER_ID"="O"."ORDER_ID")
   filter("L"."ORDER_ID"="O"."ORDER_ID")

```

EXPLAIN PLAN FOR
SELECT o.customer_id, l.unit_price * l.quantity
FROM oe.orders o, oe.order_items l
WHERE l.order_id = o.order_id;

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		665	13300	6 (17)	00:00:01	
*	1	HASH JOIN		665	13300	6 (17)	00:00:01
2	VIEW	index\$_join\$_001	105	840	3 (34)	00:00:01	

	*	3		HASH JOIN									
		4		INDEX FAST FULL SCAN		ORDER_PK		105		840		1	(0) 00:00:01
		5		INDEX FAST FULL SCAN		ORD_CUSTOMER_IX		105		840		1	(0) 00:00:01
		6		TABLE ACCESS FULL		ORDER_ITEMS		665		7980		3	(0) 00:00:01

Predicate Information (identified by operation id):

- ```

1 - access("L"."ORDER_ID"="O"."ORDER_ID")
3 - access(ROWID=ROWID)

```

10. Para la siguiente consulta:

```

SELECT *
 FROM hr.employees e, hr.departments d, hr.job_history j
 WHERE e.department_id = d.department_id
 AND e.hire_date = j.start_date;

```

- Examina su plan de ejecución.
- Examina los posibles planes resultantes de emplear *LEADING*.

#### EXPLAIN PLAN FOR

```

SELECT * FROM hr.employees e, hr.departments d, hr.job_history j
WHERE e.department_id = d.department_id AND e.hire_date = j.start_date;

```

| Id | Operation | Name                        | Rows        | Bytes | Cost (%CPU) | Time    |          |
|----|-----------|-----------------------------|-------------|-------|-------------|---------|----------|
|    | 0         | SELECT STATEMENT            |             | 11    | 1331        | 10 (20) | 00:00:01 |
| *  | 1         | HASH JOIN                   |             | 11    | 1331        | 10 (20) | 00:00:01 |
|    | 2         | TABLE ACCESS FULL           | JOB_HISTORY | 10    | 310         | 3 (0)   | 00:00:01 |
|    | 3         | MERGE JOIN                  |             | 106   | 9540        | 6 (17)  | 00:00:01 |
|    | 4         | TABLE ACCESS BY INDEX ROWID | DEPARTMENTS | 27    | 567         | 2 (0)   | 00:00:01 |
|    | 5         | INDEX FULL SCAN             | DEPT_ID_PK  | 27    |             | 1 (0)   | 00:00:01 |
| *  | 6         | SORT JOIN                   |             | 107   | 7383        | 4 (25)  | 00:00:01 |
|    | 7         | TABLE ACCESS FULL           | EMPLOYEES   | 107   | 7383        | 3 (0)   | 00:00:01 |

Predicate Information (identified by operation id):

- ```

1 - access("E"."HIRE_DATE"="J"."START_DATE")
6 - access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
   filter("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")

```

EXPLAIN PLAN FOR

SELECT /*+ LEADING(e d) */

FROM hr.employees e, hr.departments d, hr.job_history j

WHERE e.department_id = d.department_id

AND e.hire_date = j.start_date;

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		11	1331	10 (10)	00:00:01
* 1	HASH JOIN		11	1331	10 (10)	00:00:01
2	TABLE ACCESS FULL	JOB_HISTORY	10	310	3 (0)	00:00:01
* 3	HASH JOIN		106	9540	7 (15)	00:00:01
4	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01
5	TABLE ACCESS FULL	DEPARTMENTS	27	567	3 (0)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("E"."HIRE_DATE"="J"."START_DATE")
- 3 - access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")

EXPLAIN PLAN FOR

SELECT /*+ LEADING(e j) */

FROM hr.employees e, hr.departments d, hr.job_history j

WHERE e.department_id = d.department_id

AND e.hire_date = j.start_date;

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		11	1331	10 (10)	00:00:01
* 1	HASH JOIN		11	1331	10 (10)	00:00:01
* 2	HASH JOIN		11	1100	7 (15)	00:00:01
3	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01
4	TABLE ACCESS FULL	JOB_HISTORY	10	310	3 (0)	00:00:01
5	TABLE ACCESS FULL	DEPARTMENTS	27	567	3 (0)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
- 2-- access("E"."HIRE_DATE"="J"."START_DATE")

Explain plan for

SELECT /*+ LEADING(d j) */

FROM hr.employees e, hr.departments d, hr.job_history j WHERE e.department_id = d.department_id AND e.hire_date = j.start_date;

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		11	1331	32 (4)	00:00:01
* 1	HASH JOIN		11	1331	32 (4)	00:00:01
2	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01
3	MERGE JOIN CARTESIAN		270	14040	28 (0)	00:00:01
4	TABLE ACCESS FULL	DEPARTMENTS	27	567	3 (0)	00:00:01
5	BUFFER SORT		10	310	25 (0)	00:00:01
6	TABLE ACCESS FULL	JOB_HISTORY	10	310	1 (0)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID" AND
"E"."HIRE_DATE"="J"."START_DATE")

11. Para la siguiente consulta:

```
SELECT *
FROM hr.employees, hr.departments
WHERE employees.department_id = departments.department_id;
```

Evalúa si es ventajoso emplear *USE_MERGE* o *NO_USE_MERGE*.

EXPLAIN PLAN FOR

SELECT *

```
FROM hr.employees, hr.departments
WHERE employees.department_id = departments.department_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		106	9540	6 (17)	00:00:01
1	MERGE JOIN		106	9540	6 (17)	00:00:01
2	TABLE ACCESS BY INDEX ROWID	DEPARTMENTS	27	567	2 (0)	00:00:01
3	INDEX FULL SCAN	DEPT_ID_PK	27		1 (0)	00:00:01
* 4	SORT JOIN		107	7383	4 (25)	00:00:01
5	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
4 -      access("EMPLOYEES"."DEPARTMENT_ID"="DEPARTMENTS"."DEPARTMENT_ID")
      filter("EMPLOYEES"."DEPARTMENT_ID"="DEPARTMENTS"."DEPARTMENT_ID")
```

EXPLAIN PLAN FOR

SELECT /*+ USE_MERGE(employees departments) */ *

```
FROM hr.employees, hr.departments
WHERE employees.department_id = departments.department_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		106	9540	6 (17)	00:00:01
1	MERGE JOIN		106	9540	6 (17)	00:00:01
2	TABLE ACCESS BY INDEX ROWID	DEPARTMENTS	27	567	2 (0)	00:00:01
3	INDEX FULL SCAN	DEPT_ID_PK	27		1 (0)	00:00:01
* 4	SORT JOIN		107	7383	4 (25)	00:00:01
5	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
4 -      access("EMPLOYEES"."DEPARTMENT_ID"="DEPARTMENTS"."DEPARTMENT_ID")
      filter("EMPLOYEES"."DEPARTMENT_ID"="DEPARTMENTS"."DEPARTMENT_ID")
```

EXPLAIN PLAN FOR

SELECT /*+ NO_USE_MERGE(employees departments) */ * FROM

```
hr.employees, hr.departments
WHERE employees.department_id = departments.department_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		106	9540	7 (15)	00:00:01
* 1	HASH JOIN		106	9540	7 (15)	00:00:01
2	TABLE ACCESS FULL	DEPARTMENTS	27	567	3 (0)	00:00:01
3	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
1 - access("EMPLOYEES"."DEPARTMENT_ID"="DEPARTMENTS"."DEPARTMENT_ID")
```


12. En cierto código se ha encontrado la siguiente consulta:

```
SELECT /*+ NO_USE_MERGE(e d) */ *
      FROM hr.employees e, hr.departments d
      WHERE e.department_id = d.department_id
      ORDER BY d.department_id;
```

Analiza su plan de ejecución y estudia si podría mejorarse.

EXPLAIN PLAN FOR

```
SELECT /*+ NO_USE_MERGE(e d) */ *
      FROM hr.employees e, hr.departments d
      WHERE e.department_id = d.department_id
      ORDER BY d.department_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		106	9540	8 (25)	00:00:01
1	SORT ORDER BY		106	9540	8 (25)	00:00:01
* 2	HASH JOIN		106	9540	7 (15)	00:00:01
3	TABLE ACCESS FULL	DEPARTMENTS	27	567	3 (0)	00:00:01
4	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
2 - access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
```

EXPLAIN PLAN FOR

SELECT *

```
FROM hr.employees e, hr.departments d
WHERE e.department_id = d.department_id
ORDER BY d.department_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		106	9540	6 (17)	00:00:01
1	MERGE JOIN		106	9540	6 (17)	00:00:01
2	TABLE ACCESS BY INDEX ROWID	DEPARTMENTS	27	567	2 (0)	00:00:01
3	INDEX FULL SCAN	DEPT_ID_PK	27		1 (0)	00:00:01
* 4	SORT JOIN		107	7383	4 (25)	00:00:01
5	TABLE ACCESS FULL	EMPLOYEES	107	7383	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
4 -      access("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
      filter("E"."DEPARTMENT_ID"="D"."DEPARTMENT_ID")
```

13. Dada la siguiente consulta:

```
SELECT e1.last_name, e1.salary, v.avg_salary
FROM hr.employees e1,
     (SELECT department_id, avg(salary) avg_salary
      FROM hr.employees e2
      GROUP BY department_id) v
WHERE e1.department_id = v.department_id
      AND e1.salary > v.avg_salary
ORDER BY e1.last_name
```

Compara su plan de ejecución con el correspondiente a usar *MERGE* en esa misma consulta.

EXPLAIN PLAN FOR

```
SELECT e1.last_name, e1.salary, v.avg_salary FROM
hr.employees e1,
  (SELECT department_id, avg(salary) avg_salary FROM
    hr.employees e2
    GROUP BY department_id) v
WHERE e1.department_id = v.department_id AND e1.salary
      > v.avg_salary
ORDER BY e1.last_name
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		17	527	9 (34)	00:00:01
1	SORT ORDER BY		17	527	9 (34)	00:00:01
* 2	HASH JOIN		17	527	8 (25)	00:00:01
3	VIEW		11	176	4 (25)	00:00:01
4	HASH GROUP BY		11	77	4 (25)	00:00:01
5	TABLE ACCESS FULL	EMPLOYEES	107	749	3 (0)	00:00:01
6	TABLE ACCESS FULL	EMPLOYEES	107	1605	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
2 -      access("E1"."DEPARTMENT_ID"="V"."DEPARTMENT_ID")
      filter("E1"."SALARY">"V"."AVG_SALARY")
```

EXPLAIN PLAN FOR

```
SELECT /*+ MERGE(v) */ e1.last_name, e1.salary, v.avg_salary FROM hr.employees e1,
  (SELECT department_id, avg(salary) avg_salary FROM
    hr.employees e2
    GROUP BY department_id) v
WHERE e1.department_id = v.department_id AND e1.salary
      > v.avg_salary
ORDER BY e1.last_name
```

SELECT PLAN_TABLE_OUTPUT

```
FROM TABLE(DBMS_XPLAN.DISPLAY(null, null, 'ALL'));
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		165	5610	9 (34)	00:00:01
1	SORT ORDER BY		165	5610	9 (34)	00:00:01
* 2	FILTER					
3	HASH GROUP BY		165	5610	9 (34)	00:00:01
* 4	HASH JOIN		3296	109K	7 (15)	00:00:01
5	TABLE ACCESS FULL	EMPLOYEES	107	2889	3 (0)	00:00:01
6	TABLE ACCESS FULL	EMPLOYEES	107	749	3 (0)	00:00:01

Query Block Name / Object Alias (identified by operation id):

```

1 - SEL$F5BB74E1
5 - SEL$F5BB74E1 / E1@SEL$1
6 - SEL$F5BB74E1 / E2@SEL$2

```

Predicate Information (identified by operation id):

```

2 -filter("E1"."SALARY">SUM("SALARY")/COUNT("SALARY"))
4 -access("E1"."DEPARTMENT_ID"="DEPARTMENT_ID")

```

Column Projection Information (identified by operation id):

```

1 - (#keys=1) NLSSORT("E1"."LAST_NAME",'nls_sort=''SPANISH'')[210],
   "E1"."LAST_NAME"[VARCHAR2,25], "E1"."SALARY"[NUMBER,22],
   SUM("SALARY")/COUNT("SALARY")[22]
2 - "E1"."SALARY"[NUMBER,22], "E1"."LAST_NAME"[VARCHAR2,25],
   COUNT("SALARY")[22], SUM("SALARY")[22]
3 - (#keys=4) "DEPARTMENT_ID"[NUMBER,22], ROWID[ROWID,10],
   "E1"."SALARY"[NUMBER,22], "E1"."LAST_NAME"[VARCHAR2,25],
   COUNT("SALARY")[22], SUM("SALARY")[22]
4 - (#keys=1) "DEPARTMENT_ID"[NUMBER,22], ROWID[ROWID,10],
   "E1"."LAST_NAME"[VARCHAR2,25], "E1"."SALARY"[NUMBER,22],
   "SALARY"[NUMBER,22]
5 - ROWID[ROWID,10], "E1"."LAST_NAME"[VARCHAR2,25],
   "E1"."SALARY"[NUMBER,22], "E1"."DEPARTMENT_ID"[NUMBER,22]
6 - "SALARY"[NUMBER,22], "DEPARTMENT_ID"[NUMBER,22]

```

LAST_NAME	SALARY	AVG_SALARY
Abel	11000	8955,88235
Bell	4000	3475,55556
Bernstein	9500	8955,88235
...		
Vollman	6500	3475,55556
Weiss	8000	3475,55556
Zlotkey	10500	8955,88235

38 filas seleccionadas

14. Dada la siguiente consulta:

```

SELECT h.customer_id, l.unit_price * l.quantity
FROM oe.orders h, oe.order_items l
WHERE l.order_id = h.order_id;

```

Compara su plan de ejecución con el correspondiente a usar *USE_NL* en esa misma consulta.

EXPLAIN PLAN FOR

```
SELECT h.customer_id, l.unit_price * l.quantity
FROM oe.orders h, oe.order_items l
WHERE l.order_id = h.order_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		665	13300	6 (17)	00:00:01
* 1	HASH JOIN		665	13300	6 (17)	00:00:01
2	VIEW	index\$_join\$_001	105	840	3 (34)	00:00:01
* 3	HASH JOIN					
4	INDEX FAST FULL SCAN	ORDER_PK	105	840	1 (0)	00:00:01
5	INDEX FAST FULL SCAN	ORD_CUSTOMER_IX	105	840	1 (0)	00:00:01
6	TABLE ACCESS FULL	ORDER_ITEMS	665	7980	3 (0)	00:00:01

Predicate Information (identified by operation id):

```
1 - access("L"."ORDER_ID"="H"."ORDER_ID")
3 - access(ROWID=ROWID)
```

EXPLAIN PLAN FOR

```
SELECT /*+ USE_NL(l h) */ h.customer_id, l.unit_price * l.quantity
FROM oe.orders h, oe.order_items l WHERE l.order_id = h.order_id;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		665	13300	148 (2)	00:00:02
1	NESTED LOOPS		665	13300	148 (2)	00:00:02
2	VIEW	index\$_join\$_001	105	840	3 (34)	00:00:01
* 3	HASH JOIN					
4	INDEX FAST FULL SCAN	ORDER_PK	105	840	1 (0)	00:00:01
5	INDEX FAST FULL SCAN	ORD_CUSTOMER_IX	105	840	1 (0)	00:00:01
* 6	TABLE ACCESS FULL	ORDER_ITEMS	6	72	1 (0)	00:00:01

Predicate Information (identified by operation id):

```
3 - access(ROWID=ROWID)
6 - filter("L"."ORDER_ID"="H"."ORDER_ID")
```

15. Dada la siguiente consulta:

```
SELECT employee_id, department_id
FROM hr.employees
WHERE department_id > 50;
```

Estudia si su plan de ejecución puede mejorarse mediante el uso del índice *emp_department_ix*.

EXPLAIN PLAN FOR

```
SELECT employee_id, department_id
FROM hr.employees
WHERE department_id > 50;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		50	350	3 (34)	00:00:01
* 1	VIEW	index\$_join\$_001	50	350	3 (34)	00:00:01
* 2	HASH JOIN					
* 3	INDEX RANGE SCAN	EMP_DEPARTMENT_IX	50	350	2 (50)	00:00:01
4	INDEX FAST FULL SCAN	EMP_EMP_ID_PK	50	350	1 (0)	00:00:01

Predicate Information (identified by operation id):

```
-----
1 - filter("DEPARTMENT_ID">50)
2 - access(ROWID=ROWID)
3 - access("DEPARTMENT_ID">50)
```

EXPLAIN PLAN FOR

```
SELECT /*+ INDEX (employees emp_department_ix)*/ employee_id, department_id FROM
  hr.employees
WHERE department_id > 50;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		50	350	6 (0)	00:00:01
1	TABLE ACCESS BY INDEX ROWID	EMPLOYEES	50	350	6 (0)	00:00:01
* 2	INDEX RANGE SCAN	EMP_DEPARTMENT_IX	50		1 (0)	00:00:01

Predicate Information (identified by operation id):

```
-----
2 - access("DEPARTMENT_ID">50)
```

16. En cierto código se ha encontrado la siguiente consulta:

```
SELECT /*+ NO_INDEX(employees emp_empid) */ employee_id
FROM hr.employees
WHERE employee_id > 200;
```

Analiza su plan de ejecución e indica si convendría realizar algún cambio.

EXPLAIN PLAN FOR

```
SELECT /*+ NO_INDEX(employees emp_empid) */ employee_id
FROM hr.employees
WHERE employee_id > 200;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		6	24	1 (0)	00:00:01
* 1	INDEX RANGE SCAN	EMP_EMP_ID_PK	6	24	1 (0)	00:00:01

Predicate Information (identified by operation id):

```
-----
1 - access("EMPLOYEE_ID">200)
```

EXPLAIN PLAN FOR

```
SELECT employee_id
FROM hr.employees
WHERE employee_id > 200;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		6	24	1 (0)	00:00:01
* 1	INDEX RANGE SCAN	EMP_EMP_ID_PK	6	24	1 (0)	00:00:01

Predicate Information (identified by operation id):

1 - access("EMPLOYEE_ID">200)

EXPLAIN PLAN FOR

**SELECT /*+ NO_INDEX(employees emp_emp_id_pk) */ employee_id
FROM hr.employees
WHERE employee_id > 200;**

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		6	24	3 (0)	00:00:01
* 1	TABLE ACCESS FULL	EMPLOYEES	6	24	3 (0)	00:00:01

Predicate Information (identified by operation id):

1 - filter("EMPLOYEE_ID">200)